

Precise Divide-By-Zero Detection with Affirmative Evidence

Yiyuan Guo
The Hong Kong University of Science
and Technology
Hong Kong, China
yguoaz@cse.ust.hk

Jinguo Zhou
Ant Group
China
jinguo.zjg@antgroup.com

Peisen Yao
The Hong Kong University of Science
and Technology
Hong Kong, China
pyao@cse.ust.hk

Qingkai Shi
Ant Group
China
qingkai.sqk@antgroup.com

Charles Zhang
The Hong Kong University of Science
and Technology
Hong Kong, China
charlesz@cse.ust.hk

ABSTRACT

The static detection of divide-by-zero, a common programming error, is particularly prone to false positives because conventional static analysis reports a divide-by-zero bug whenever it cannot prove the safety property — the divisor variable is not zero in all executions. When reasoning the program semantics over a large number of under-constrained variables, conventional static analyses significantly loose the bounds of divisor variables, which easily fails the safety proof and leads to a massive number of false positives. We propose a static analysis to detect divide-by-zero bugs taking additional evidence for under-constrained variables into consideration. Based on an extensive empirical study of known divide-by-zero bugs, we no longer arbitrarily report a bug once the safety verification fails. Instead, we actively look for affirmative evidences, namely source evidence and bound evidence, that imply a high possibility of the bug to be triggerable at runtime. When applying our tool `WIT` to the real-world software such as the Linux kernel, we have found 72 new divide-by-zero bugs with a low false positive rate of 22%.

CCS CONCEPTS

• **Software and its engineering** → **Software verification and validation.**

KEYWORDS

Static program analysis, bug detection, divide-by-zero.

ACM Reference Format:

Yiyuan Guo, Jinguo Zhou, Peisen Yao, Qingkai Shi, and Charles Zhang. 2022. Precise Divide-By-Zero Detection with Affirmative Evidence. In *44th International Conference on Software Engineering (ICSE '22)*, May 21–29, 2022, Pittsburgh, PA, USA. ACM, New York, NY, USA, 12 pages. <https://doi.org/10.1145/3510003.3510066>

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](https://permissions.acm.org).
ICSE '22, May 21–29, 2022, Pittsburgh, PA, USA

© 2022 Association for Computing Machinery.
ACM ISBN 978-1-4503-9221-1/22/05...\$15.00
<https://doi.org/10.1145/3510003.3510066>

1 INTRODUCTION

Divide-by-zero is one of the most common programming errors that can lead to undefined runtime behaviors. Over the past five years, there are more than 90 security vulnerabilities related to divide-by-zero documented in the Common Vulnerabilities and Exposures (CVE) database. Static analysis is an important approach to combat this problem. However, its high false positive rate remains a major obstacle for adoption. In our experiments, we observe false positive rates of over 70% in representative static analyzers.

To understand the reason behind these high false positive rates, we first examine how static analysis reports a potential divide-by-zero bug. The central question is how to decide if the divisor variable may equal zero in some executions. More specifically, a divide-by-zero happens when

- C1. There exists a feasible execution path P reaching the division instruction (Let pc denote the path condition for P).
- C2. The divisor variable v may equal zero when P reaches the division instruction (Let $Q \stackrel{\text{def}}{=} (v = 0)$ denote the error condition for triggering a divide-by-zero bug).

Conventionally, static analysis reports a divide-by-zero bug if $pc \wedge Q$ is satisfiable, i.e., under the condition pc , the divisor variable v cannot be proved to be non-zero. However, we observe that this scheme often leads to a high false positive rate in divide-by-zero detection since the under-constrainedness of static analysis makes it easy to satisfy the query $pc \wedge Q$. Many variables are under-constrained in the static analysis [13]:

- The program under analysis can be an open program. For example, the analysis often sees a module by interface only, of which the variables are under-constrained [11, 15].
- The inherent limitations of static analysis introduce the under-constrained variables to represent imprecision in the analysis, e.g., the runtime state of the operating system and the presence of unmodeled code constructs, such as inline assembly and many others [13, 14].

To improve the precision, our key insight is that, instead of reporting errors whenever safety verification fails, we can attack the divide-by-zero detection problem from a different angle by actively finding affirmative evidence for triggering the bug. Such evidence serves as the extra information on the under-constrained variables, thus contributes to improving the precision.

```

1 struct control_panel {
2     int x; int y; int z;
3     int flag; int distance; ...
4 };
5
6 void top_fun() { // not tracked by the analysis when analyzing move
7     ...
8     if (dx != 0 && dy != 0)
9         // calls move with non-zero values of dx and dy.
10        move(dx, dy, d, ctr);
11 }
12
13 void move (int dx, int dy, int d, control_panel *ctr) {
14     int diff;
15     if (dx >= dy)
16         diff = 1 + dx - dy;
17     else
18         diff = -1;
19
20     ctr->z += get_step(ctr, diff);
21     if (dx >= d || dy >= d)
22         ctr->setFlag(LONG_MOVE);
23
24     int dz = 2*dx - (dy+d);
25     ctr->x += get_step(ctr, dx);
26     ctr->y += get_step(ctr, dy);
27     ctr->z += get_step(ctr, dz);
28 }
29
30 int get_step(control_panel *ctr, int step_size) {
31     return ctr->distance / step_size;
32 }

```

Figure 1: A motivating example.

We use the motivating example in Figure 1 to illustrate the idea. In Figure 1, the function `move` updates the control panel `ctr` by adding values to its (x, y, z) coordinates computed in the function `get_step` through a division operation. Also, we assume that the caller function `top_fun` for the function `move` is external to the module under analysis and, thus, not tracked by the analysis. Now consider how static analysis can detect potential divide-by-zero errors in this program.

Line 20 cannot trigger a divide-by-zero because the used divisor `diff` is either -1 or no less than 1 . A path-sensitive static analysis correctly verifies its safety since the two queries $dx \geq dy \wedge 1 + dx - dy = 0$ and $dx < dy \wedge -1 = 0$ associated with the two paths across Lines 15-18 are both unsatisfiable.

However, the path-sensitive static analysis still encounters the precision problem. In our example, the divide-by-zero errors are reported for the three calls to `get_step` at Line 25, 26, and 27 because the queries $dx = 0$, $dy = 0$, and $dz = 0$ are all satisfiable. However, since `top_fun` only calls `move` with non-zero values for `dx` and `dy` (Line 10), the reports at Lines 25-26 are false positives. This fact is unknown to the static analysis because it fails to trace back to the origin of the arguments of `move` in function `top_fun`.

On the other hand, we notice that the divide-by-zero report at Line 27 has high confidence to be true since we can find evidence to explain it. Specifically, Line 21 of Figure 1 explicitly compares `dx` and `dy` with `d`, suggesting that the programmer has beliefs [16] that they may be equal to `d`. If such beliefs actually hold, `dz` must equal zero at Line 24, i.e., $dz = 2 \cdot dx - (dy + d) = 2 \cdot d - (d + d) = 0$, leading to a plausible divide-by-zero report for Line 27.

How to report the high confidence divide-by-zero bug at Line 27 instead of the false positives at Lines 25-26? Note that all of the

three divisions cannot be proved safe by the static analysis, regardless of being path-sensitive or not. However, we have identified the evidence based on analyzing the programmer's beliefs for the bug report at Line 27, which leads to its high confidence to occur. Specifically, in this work, we identify two categories of evidence:

- (a) **Source evidence:** The fact that an explicit source of “bad” value is assigned to a variable v . The source can be either a direct assignment of zero (e.g., the assignment $v := 0$) or takes the value of some tainted input (e.g., $v := \text{atoi}(\text{argv}[1])$).
- (b) **Bound evidence:** The equality fact $v_1 = v_2$ generated from a bound checking statement in the program that compares v_1 and v_2 . Statements like Line 21 of Figure 1 that check numerical bounds of variables can convey the important information on the possible values of the checked variables: a comparison between v_1 and v_2 suggests that the programmer may believe that v_1 can equal v_2 .

With the generated evidence, our method reports a divide-by-zero bug by adapting the requirements C1 and C2 as follows.

- C1*. C1 and C2 hold: $pc \wedge Q$ is satisfiable, where pc denotes the path condition for reaching the division instruction, and $Q \equiv (v = 0)$ is the error condition for divide-by-zero.
- C2*. One of the following conditions holds:
 - The divisor variable v has the source evidence.
 - There exists a set E of bound evidence consistent with pc such that v **must** be zero under E .

In other words, we seek to detect a fraction of divide-by-zero bugs with high confidence by finding the affirmative evidence to explain its occurrence. For example, in Figure 1, `dz` at Line 27 is the only variable that meets these requirements: $dz = 0$ **must** hold if the evidence $dx = d$ and $dy = d$ hold. Hence, our approach will only report one divide-by-zero bug at Line 27.

In this paper, we propose WIT, a framework for the precise detection of divide-by-zero with the affirmative evidence. First, to understand the applicability of our definition of evidence and the criteria C1* - C2* for detecting bugs, we perform an empirical study on existing CVEs related to divide-by-zero bugs and investigate if they can be detected by finding a set of supporting evidence. The result shows that 74% of the studied divide-by-zero bugs have the supporting evidence that explains its occurrence, showing the generality of our intuition. To capture the evidence-based reasoning in achieving the precise divide-by-zero detection, we design a static analysis algorithm to perform the evidence-based inference, adhering to the criteria C1*-C2* for reporting bugs. The algorithm generates evidence directly from certain code patterns and propagates the generated evidence path-sensitively. To improve its efficiency, we utilize a symbolic domain to compactly encode the possible numerical values for variables and apply the data dependence analysis [17] to prune irrelevant execution paths, scaling to million lines of code.

In summary, we make the following contributions in this paper:

- The insight for improving the precision of divide-by-zero detection by finding the affirmative evidence to trigger the bug.
- An empirical study of CVEs related to divide-by-zero bugs. We investigate and classify these existing divide-by-zero

bugs and show that many of them can be effectively detected based on evidence.

- A formalization of the insight in a semantic framework and an algorithm for finding high confidence divide-by-zero errors through the evidence propagation.
- An implementation and extensive evaluation of the divide-by-zero checker. We demonstrate that it is both precise and efficient, uncovering 72 divide-by-zero bugs (14 of which are confirmed by the developers) in large codebases such as the Linux kernel with a low false positive rate of 22%.

2 EMPIRICAL STUDY ON DIVIDE-BY-ZERO BUGS

To further understand the applicability of the evidence-based divide-by-zero detection method, we perform an empirical study on CVEs caused by divide-by-zero bugs. Through the empirical study, we aim to answer the following research questions:

- **RQ1:** How often can we find the evidence for the divide-by-zero bugs? In other words, how often do our bug detection criteria C1* - C2* apply to the existing divide-by-zero bugs?
- **RQ2:** What is the distribution for the two kinds of evidence defined in § 1?

2.1 Data Collection

We search for the keywords “divide-by-zero” and “divide by zero” in the CVE database and examine the CVEs starting from the year 2011. There are 123 CVEs in total that are caused by the divide-by-zero bugs. We exclude the bugs with no source code or stack traces (16), or unable to understand without a deep knowledge of the system (12), or sharing the same root with other bugs (8). Thus, we are left with 87 CVEs to study.

2.2 Classification Criteria

Since our goal is to study the problem in general without being tied to a specific analysis algorithm or target system, we have adopted the following criteria to mimic the reasoning process of a static analyzer:

- (1) Starting from the crash site of the bug (i.e., the division instruction with zero divisor), we manually examine a backward slice \mathcal{B} of some fixed size. In our experiment, we examine backwards at most 10 call frames from the involved division. This is reasonable as lengthy bug traces output by static analyzers take a non-trivial amount of time for users to triage [1, 4], thus should be avoided by practical tools.
- (2) Based on our manual inspection of \mathcal{B} , we classify the divide-by-zero bugs based on the evidence found:
 - (a) *Class Src:* Source evidence is found. An explicit source of zero value for the divisor variable exists in \mathcal{B} . The source is either a direct assignment of zero value to the variable or of a tainted value from the input.
 - (b) *Class Bd:* Bound evidence is found. We can find a set of evidence from the bound checking statements in \mathcal{B} that guarantee the divisor variable to be zero. Specifically, for any branching statement involving comparisons:

Table 1: Classification of 87 CVEs caused by divide-by-zero bugs. Class Src is bugs with source evidence (further classified into constant source and taint source), Class Bd is bugs with bound evidence, and Class Un is bugs with no evidence.

Class Src		Class Bd	Class Src \cap Class Bd	Class Un
const	taint			
19	28	21	4	23

if $(x \text{ cmp } y)$ (cmp is $<$, \leq , etc.), we consider the fact $x=y$ as the bound evidence likely to be true.

- (c) *Class Un:* No evidence is found (i.e., unknown). We can neither find source evidence nor bound evidence in \mathcal{B} . For example, CVE-2018-19628 is marked as unknown because it requires a deep context of 14 call frames to understand the root cause, exceeding the code range of \mathcal{B} .

Conventional static analysis methods detect all three classes of divide-by-zero bugs. As illustrated in § 1, they are likely to incur massive false positives. Our method detects bugs of Class Src and Class Bd since they are the classes that satisfy the requirement C1* - C2*, aiming for a fraction of high confidence divide-by-zero bugs with possible false negatives. The classification attempts to study how much can be covered by our method.

2.3 Study Result

Bug Classification. Table 1 shows the classification result of the 87 CVEs caused by divide-by-zero bugs. For Class Src bugs with explicit sources, we further divide them into two groups: those with constant zero value as the source (the column labeled with “const”) and those with the tainted input as the source (the column labeled with “taint”). The bugs with both the source evidence and the bound evidence are counted in the column labeled “Class Src \cap Class Bd”.

From the data we gather, we answer the two research questions empirically as follows:

Answer to RQ1: We can find evidence for a large proportion ($\frac{Src+Bd-Src \cap Bd}{87} = 74\%$) of divide-by-zero bugs. Thus, our bug detection criteria C1* - C2* are widely applicable.

Answer to RQ2: For the divide-by-zero bugs that we can find evidence, 73% of them have source evidence, 33% of them have bound evidence, and 6% of them have both evidence.

Therefore, we conclude that finding evidence to facilitate the precise divide-by-zero detection indeed works in real scenarios, covering a large proportion of interesting bugs. Also, both source evidence and bound evidence are useful for bug detection in practice.

Case Study of Bugs with Evidence. To give an intuitive understanding of the bugs detected based on evidence, we illustrate some example CVEs from Class Src and Bd.

Figure 2a shows a Class Src bug. The buffer `hdrbuf` is tainted by the call to `ReadOK` (Line 4), allowing it to parse the variable width to zero value, leading to a divide-by-zero bug at Line 9.

Figure 2b shows a Class Bd bug. At Line 5 of the code in Figure 2b, the programmer explicitly checks the value of variable v

```

1 void start_input_gif(j_compress_ptr cinfo, jpeg_source_ptr sinfo) {
2   U_CHAR hdrbuf[10];
3   // ReadOK is a wrapper of fread
4   if (!ReadOK(source->pub.input_file, hdrbuf, 6))
5     EREXIT(cinfo, JERR_GIF_NOT);
6   width = LM_to_uint(hdrbuf, 0);
7   height = LM_to_uint(hdrbuf, 2);
8   // use the second arg as divisor
9   DIVOP(cinfo, width * NUMCOLORS);
10 }

```

(a) Example #1: CVE-2021-20205

```

1 int ApplyEvaluateOperator(RandomInfo *r, Quantum p, Operator op, int v) {
2   ...
3   switch(op) {
4     case DivideEvaluateOperator:
5       result = p / (v == 0 ? 1 : v);
6       break;
7     case GaussianNoiseEvaluateOperator:
8       result = GenerateDifferentialNoise(r, p, GaussianNoise, v)
9   }
10 }
11 int GenerateDifferentialNoise(RandomInfo *r, Quantum p, NoiseType ty, int v)
12 {
13   ...
14   noise = (QuantumRange * i / v);
15 }

```

(b) Example #2: CVE-2021-20176

Figure 2: Example of bugs with evidence.

against 0. We can propagate the checked bound as evidence to Line 8, uncovering a divide-by-zero bug inside the called function `GenerateDifferentialNoise` at Line 13. This bug is due to the fact that the programmer has missed the check at Line 8, which is successfully captured by the generated bound evidence.

Implications of our Findings. From the findings of our empirical study, we conclude that:

- (1) Many real divide-by-zero bugs do have the supporting evidence. In addition, both the source evidence and the bound evidence are useful for understanding existing divide-by-zero bugs. Thus, criteria $C1^*$ - $C2^*$ for the divide-by-zero detection are applicable for a large proportion of divide-by-zero bugs.
- (2) Static analyzers can be designed to discover and reason about the evidence to achieve precise detection results, aiming for the high confidence bugs of Class Src and Class Bd.

3 OVERVIEW

In this section, we give an overview of our evidence-based divide-by-zero detection method, using the motivating example in Figure 1.

Existing Approaches. Before presenting our approach, we first briefly discuss two conventional techniques for divide-by-zero detection and their drawbacks, namely the numerical abstract interpretation [8, 9, 26] and the under-constrained symbolic execution [28].

If we apply the polyhedra abstract domain [9] that is an expressive domain capable of reasoning about the linear relations among variables, the branch statements from Line 15 to Line 18 in Figure 1

will produce the following abstractions along the two paths:

$$Abs_1 : \{dx - dy \geq 0, dx - dy - diff + 1 = 0\}$$

$$Abs_2 : \{dx - dy \leq 0, diff + 1 = 0\}$$

Since the two abstractions are joined to produce a sound over-approximation when paths merge in abstract interpretation, a severe precision loss can occur. Specifically, since $Abs_1 \sqcup Abs_2 = \{diff \geq -1\}$, the analysis concludes that `diff` may be equal to zero and generates a false positive divide-by-zero bug for the call at Line 20. Thus, lacking path sensitivity, the numerical abstract interpretation can be imprecise for divide-by-zero detection.

On the other hand, the under-constrained symbolic execution performs the path-by-path exhaustive exploration, starting from the tested function move (recall we assume that the analysis does not track the function `top_fun`). The analysis records the two paths in the function move reaching Line 20 as the following execution states:

$$s_1 = (dx \geq dy) \wedge (diff = 1 + dx - dy)$$

$$s_2 = (dx < dy) \wedge (diff = -1)$$

With the path-sensitive state encoding, the under-constrained symbolic execution successfully proves that divide-by-zero cannot happen for Line 20. However, as illustrated in § 1, due to the presence of the under-constrained variables, we conclude that all three calls at Lines 25-27 of Figure 1 can trigger divide-by-zero with two false positives. This is because all the execution paths reaching these lines have satisfiable constraints for triggering divide-by-zero. Thus, without finding the affirmative evidence, even a path-sensitive technique for divide-by-zero detection can be imprecise when faced with under-constrained variables.

To summarize, conventional approaches to divide-by-zero detection suffer from imprecision problems because they may lack path sensitivity or do not attempt to find evidence for potential bugs.

Our Approach. As illustrated in § 1, our work uses the criteria $C1^*$ - $C2^*$ to precisely report potential bugs. We argue that both path sensitivity of the analysis and finding affirmative evidence are crucial in achieving good precision. Thus, we seek to track the evidence path-sensitively to achieve the precise divide-by-zero detection.

For this purpose, the analysis needs to reason about the numerical computations path-sensitively and fuse the evidence in the analysis process. We propose to apply a dedicated symbolic domain Γ to track the numerical computation in the analysis path-sensitively and utilize the evidence to refine the analysis result. Specifically, Γ represents the abstract state for a variable v as a guarded symbolic value set $\Gamma(v) = \{(va_1, c_1), \dots, (va_k, c_k)\}$, meaning that $v == va_i$ may hold under the condition c_i . Note that it is nontrivial to compute Γ efficiently, and we defer the details of the algorithm to § 4. For the example in Figure 1, our analysis deduces that

$$(1 + dx - dy, dx \geq dy), (-1, dx < dy) \in \Gamma(diff)$$

$$(2 \times dx - (dy + d), true) \in \Gamma(dz)$$

The analysis successfully infers that no divide-by-zero can occur for Line 20 based on the path-sensitive representation of $\Gamma(diff)$. However, it still reports two spurious divide-by-zero bugs for Lines 25-26,

similar to the under-constrained symbolic execution approach, as the variables dx , dy , and d are under-constrained in Γ .

Thus, to further improve the precision, our analysis attempts to deduce the evidence for those under-constrained variables. First, it attempts to find the **source evidence** for dx , dy , and d by asking “Are these variables tainted?”. The tracking of the source evidence can be done by the taint analysis [31], following the value that originates from a bad source and checking where it flows to. In this example, no such source evidence exists. On the other hand, the analysis tries to find the **bound evidence** to constrain the under-constrained variables. As introduced in § 1, we generate the bound evidence $v_1 = v_2$ from the bound checking statements that compare v_1 with v_2 , which represent the possible beliefs the programmers may have. For example, in Figure 1, Lines 15 compares dx with dy , while Line 21 compares dx with d and dy with d , causing the analysis to generate the following bound evidence:

$$dx = dy, dx = d, dy = d$$

How can we take advantage of the generated bound evidence to improve the precision of divide-by-zero detection? We treat the bound evidence as the additional constraints and propagate it to update the representation of Γ , such that the additional constraints are enforced. For the variable dz , from the bound evidence above and the guarded value $(2 \times dx - (dy + d), true)$ for dz , we deduce that

$$(2 \times a - (b + c), true) \in \Gamma(dz), a, b, c \in \{dx, dy, d\}$$

Apparently, $0 \in \Gamma(dz)$ and thus our analysis successfully reports Line 27 of Figure 1 as a divide-by-zero bug. Meanwhile, since no evidence is inferred for dx or dy to be zero (i.e., $0 \notin \Gamma(dx), 0 \notin \Gamma(dy)$), it avoids generating spurious reports at Lines 25-26.

In summary, our analysis encodes the possible values for the variables in the program path-sensitively using a symbolic domain. During the numerical inference process, it attempts to (1) directly find the source evidence and (2) propagate the generated bound evidence to refine the analysis result. Since we apply standard methods in tracking the source evidence, we mainly focus on utilizing the bound evidence in our method. There are two major challenges:

- How to efficiently compute the guarded symbolic value set $\Gamma(v)$ for variable v ? A naive approach that exhaustively enumerates all execution paths can be too expensive, wasting time exploring irrelevant paths.
- How to propagate bound evidence to refine the symbolic analysis result?

To address the first challenge, we utilize the data dependence analysis to slice away the irrelevant control flow paths for improving the efficiency. For the second challenge, we encode the bound evidence as the additional constraints for Γ and enforce these constraints when computing the guarded symbolic value sets.

4 METHODOLOGY

4.1 Preliminary Definitions

We formulate our analysis using a simple imperative programming language defined in Figure 3. The language is assumed in the static single assignment form [10] in which each variable has a unique definition, and we denote the SSA phi function by ϕ . We use $v@l$ to

```

Program P ::= fun+
Function fun ::= fun : (v1, ..., vn) → r
               {s}
Statement s ::= s1; s2 | v := e
              | if (v1 cmp v2) s2 else s3, cmp ∈ {≤, =}
              | v := g(a1, ..., an) | v := ϕ(v1, ..., vn)
Expression e ::= v | c | tainted | e1 op e2, op ∈ {+, −, ×, ÷}

```

Figure 3: A simple programming language.




Figure 4: Augmented data dependence graph for Figure 1.

denote a variable v defined at the program location l . The language has standard semantics, and we omit a formal definition for brevity.

As mentioned in § 3, we utilize data dependence analysis to compute Γ efficiently. Specifically, the analysis operates over a sparse representation of the program called the augmented data dependence graph, defined as follows:

Definition 4.1. An augmented data dependence graph \mathcal{G} is a directed graph $\mathcal{G} = (V, E, L_V, L_E)$ where:

- (1) $V = V_C \cup V_N \cup V_O \cup V_t$ is the set of nodes. V_C are constant values, V_N are variable definitions, and V_O are nodes corresponding to an arithmetic operation $v_1 \text{ op } v_2$, where $v_1, v_2 \in V \setminus V_O$. V_t is the set of tainted input sources.
- (2) $E \subseteq V \times V$ is the set of edges representing data dependence relations such that $e = (v_1, v_2) \in E$ when v_1 is used to define v_2 . L_E labels any $e \in E$ with a condition $cond$ under which the value flow can happen.
- (3) L_V labels each node $n \in V_N$ with $(v@l, cond)$. $cond$ is the condition for v 's definition (i.e., for some execution to reach Line l). L_V also labels nodes in V_O with its arithmetic expression $v_1 \text{ op } v_2$ and nodes in V_C with its constant c .

Example 4.1. Figure 4 shows the augmented data dependence graph for the program in Figure 1. An arrow $v_1 \rightarrow v_2$ indicates that v_1 is used to define v_2 (i.e., $(v_1, v_2) \in E$), and each arrow is labeled with a path condition under which the flow of value can happen (omitted in Figure 4 if it is *true*). Notice that arithmetic operations are also compactly encoded on the graph by introducing temporary variables t_1 , t_2 , and t_3 to represent intermediate computation results.

The sparse graph representation is used to track relevant data and control dependencies of the concerned variable while skipping irrelevant statements [17]. It has been previously shown effective in detecting null pointer dereference [3, 29], use after free [29], and memory leak [6, 30]. We take inspiration from these works to first utilize data dependence analysis in finding divide-by-zero bugs.

As illustrated in § 3, to achieve precise divide-by-zero detection, our analysis needs to track numerical computations path-sensitively

and utilize affirmative evidence to find high confidence bugs. For tracking numerical computations, we apply a symbolic domain Γ , formally defined as follows:

Definition 4.2. The guarded symbolic value set domain is a mapping $\Gamma \in V \rightarrow \mathcal{P}(Val \times Cond)$, where:

$$Val \stackrel{\text{def}}{=} \{c, \hat{v} | c \in \mathbb{Z}, v \in V\} \cup \{va_1 \text{ op } va_2 \mid va_1, va_2 \in Val\}$$

$Cond \stackrel{\text{def}}{=} \text{the set of path conditions}$

We use \hat{v} to denote an unknown symbolic value for a node v .

Γ encodes the symbolic values for variables path-sensitively. It maps a node $v \in \mathcal{G}$ to guarded value pairs of the form (va, c) , where $va \in Val$ is a symbolic expression and c is the condition under which v may have the value va . The symbolic expression in Val is either a basic term (i.e., constant or unknown symbolic value) or a binary operation involving other symbolic expressions.

At a high level, our analysis first constructs the augmented data dependence graph \mathcal{G} utilizing the existing method [29]. We then compute the guarded symbolic value set $\Gamma(v)$ for a divisor variable v on-demand. Finally, we use the result $\Gamma(v)$ in detecting divide-by-zero bugs. Specifically, we compute $\Gamma(v)$ by building and resolving a system of constraints for Γ . To improve precision, our analysis is path-sensitive and evidence aware, encoding bound evidence as additional constraints. To remain efficient, we utilize data dependence relations in constructing and resolving constraints.

4.2 Evidence-based Symbolic Analysis

Given a program P and a variable v , our evidence-based symbolic analysis shown in Algorithm 1 computes the guarded symbolic value set for v . We build the augmented data dependence graph \mathcal{G} for P using standard methods (Line 2), generate a system of constraints $Cons(\Gamma)$ on-demand (Line 4), and solve the constraints to obtain the guarded symbolic value set for v (Lines 5-6).

§ 4.2.1 demonstrates the process of on-demand constraints generation for a node $n \in \mathcal{G}$, which is the node corresponding to a given variable v . We also generate bound evidence to produce additional constraints, further refining the result and improving the precision. § 4.2.2 discusses the procedure *solve* for resolving constraints and obtaining the solution Γ , which maps from variables to their guarded symbolic value sets.

4.2.1 Symbolic Constraints Generation. Our analysis generates constraints for Γ from the program's augmented data dependence graph \mathcal{G} . Before presenting the rules for constraint generation, we first define some operators for Γ (cf. Definition 4.2):

Definition 4.3. Operators and helper functions definitions for Γ :

- (1) The binary operator $\overline{\text{op}}$ on $\mathcal{P}(Val \times Cond)$ is defined as:
 $GV_1 \overline{\text{op}} GV_2 = \{\text{simplify}(va_1 \text{ op } va_2), c_1 \wedge c_2 \mid (va_1, c_1) \in GV_1, (va_2, c_2) \in GV_2\}$
- (2) The logical and operation \wedge is extended to add an additional condition to a guarded value set $S \in \mathcal{P}(Val \times Cond)$: $cond \wedge S = \{(va, c \wedge cond) \mid (va, c) \in S\}$.
- (3) $\dot{\cup}$ unions two guarded value sets $S_1, S_2 \in \mathcal{P}(Val \times Cond)$ and combines guarding conditions for the same value: $S_1 \dot{\cup} S_2 = \{(va, c_1 \wedge c_2) \mid (va, c_1) \in S_1, (va, c_2) \in S_2\} \cup (S_1 \setminus S_2) \cup (S_2 \setminus S_1)$.

Algorithm 1: Demand-driven evidence-based symbolic analysis.

Input: A program P and a cared variable v .

Output: The guarded symbolic value set for v .

```

1 def symbolicAnalysis( $P, v$ ):
2    $\mathcal{G}(V, E) \leftarrow \text{buildGraph}(P)$ 
3    $n \leftarrow \text{node in } \mathcal{G} \text{ corresponding to } v$ 
4    $Cons(\Gamma) \leftarrow \text{genConstraints}(\mathcal{G}, n)$ 
5    $\Gamma \leftarrow \text{solve}(Cons(\Gamma))$ 
6   return  $\Gamma(n)$ 
7 def genConstraints( $\mathcal{G}, n$ ):
8    $Cons(\Gamma) \leftarrow \emptyset$ 
9    $Seq = \{v_i\} \leftarrow \text{backward dfs order of } \mathcal{G} \text{ starting from } n$ 
10  foreach  $v \in \text{Reverse}(Seq)$  do
11    Apply the rules in Figure 5 and add the generated
    constraints to  $Cons(\Gamma)$ 
12  end
13  return  $Cons(\Gamma)$ 
14 def solve( $Cons(\Gamma)$ ):
15    $\Gamma \leftarrow \lambda v. \emptyset$ 
16   while  $\Gamma$  has changed do
17     foreach  $(C : \Gamma(o) \leftrightarrow S) \in Cons(\Gamma)$  do
18        $S' \leftarrow \Gamma(o) \dot{\cup} \text{eval}(S, \Gamma)$ 
19        $\Gamma \leftarrow \Gamma[o \mapsto S']$ 
20     end
21  end
22  return  $\Gamma$ 

```

Figure 5 lists the inference rules for generating constraints for Γ based on \mathcal{G} . We use $S_1 \leftrightarrow S_2$ to denote a weak update to the guarded value set S_1 (i.e., $S_1 = S_1 \dot{\cup} S_2$), and \xrightarrow{cond} to denote data dependence edges (c.f. Definition 4.1). The first five rules in Figure 5 translate the data dependence relations in \mathcal{G} to constraints for Γ . Rule *init-var* and Rule *init-op* indicate that every node $n \in V_n \cup V_o$ is associated with an unknown symbolic value \hat{n} , meaning that the value for n is initially unknown. On the other hand, constant nodes have fixed values (Rule *init-cst*). For node encoding operations, the operator $\overline{\text{op}}$ is applied to the guarded symbolic value sets for the incoming nodes (Rule *operation*)¹. For a variable node $n \in V_n$, we aggregate the guarded symbolic value sets for all the nodes that n is data dependent on, and track the appropriate path conditions (Rule *variable*). The *trans* function is used for this reason and defined as:

$$\text{trans}([\Gamma(op_1), c_1], \dots, [\Gamma(op_k), c_k]) \stackrel{\text{def}}{=} \bigcup_{i=1}^k (c_i \wedge \Gamma(op_i))$$

Finally, the condition associated with n in $L_V(n)$ is added to $\Gamma(n)$.

Example 4.2. For \mathcal{G} shown in Figure 4, we can apply the above inference rules to build the following constraints (the constraints

¹For brevity, the rule only lists the case when both operand nodes are from V_n . Other cases can be transformed to it by introducing temporary variables.

$$\begin{array}{c}
\text{init-var} \frac{n \in V_n, L_V(n) = (v, \text{cond})}{\Gamma(n) \leftarrow \{(\hat{n}, \text{cond})\}} \\
\\
\text{init-op} \frac{n \in V_o}{\Gamma(n) \leftarrow \{(\hat{n}, \text{true})\}} \\
\\
\text{init-cst} \frac{n \in V_c, L_V(n) = \text{cst}}{\Gamma(n) \leftarrow \{(\text{cst}, \text{true})\}} \\
\\
\text{operation} \frac{\begin{array}{l} n \in V_o, L_V(n) = v_1 \text{ op } v_2 \\ op_1 \rightarrow n, L_V(op_1) = (v_1, c_1) \\ op_2 \rightarrow n, L_V(op_2) = (v_2, c_2) \end{array}}{\Gamma(n) \leftarrow \Gamma(op_1) \widetilde{\text{op}} \Gamma(op_2)} \\
\\
\text{variable} \frac{\begin{array}{l} op_1 \xrightarrow{\text{cond}_1} n, \dots, op_k \xrightarrow{\text{cond}_k} n \\ n \in V_n, L_V(n) = (v@l, \text{cond}) \\ vas = \text{trans}([\Gamma(op_1, \text{cond}_1), \dots, (\Gamma(op_k), \text{cond}_k)]) \end{array}}{\Gamma(n) \leftarrow \text{cond} \wedge vas} \\
\\
\text{bound-evi} \frac{v_1 \text{ cmp } v_2 \text{ occurs in the program}}{\Gamma(v_i) \leftarrow \Gamma(v_1) \dot{\cup} \Gamma(v_2), i \in \{1, 2\}}
\end{array}$$

Figure 5: Rules for generating constraints from \mathcal{G} . $S_1 \leftrightarrow S_2$ abbreviates $S_1 = S_1 \dot{\cup} S_2$, and $i \xrightarrow{\text{cond}} j$ abbreviates $e = (i, j) \in E, L_E(e) = \text{cond}$.

for temporary variables are left out for brevity):

$$\Gamma(\text{step_size}) \leftarrow \Gamma(\text{diff}) \dot{\cup} \Gamma(dx) \dot{\cup} \Gamma(dy) \dot{\cup} \Gamma(dz) \quad (1)$$

$$\Gamma(\text{diff}) \leftarrow \text{trans}([\Gamma(-1), dx < dy, \quad (2)$$

$$(\Gamma(1) \tilde{\wedge} (\Gamma(dx) \tilde{\wedge} \Gamma(dy)), dx \geq dy)]) \quad (3)$$

$$\Gamma(dz) \leftarrow \Gamma(2) \tilde{\wedge} \Gamma(dx) \tilde{\wedge} (\Gamma(dy) \tilde{\wedge} \Gamma(d)) \quad (4)$$

$$\Gamma(v) \leftarrow \{(\hat{v}, \text{true})\}, v \in \{\text{step_size}, dx, dz, \text{diff}, dy, d\} \quad (5)$$

Since we build the constraints from the augmented data dependence graph shown in Figure 4, only relevant path conditions are encoded in the constraint representation. For instance, no branch condition in Figure 1 can affect the value of dz , i.e., $dz = 2 \times dx - (dy + d)$ holds on all eight execution paths. The above constraint (4) effectively captures this fact, indicating that $\Gamma(dz)$ is dependent on $\Gamma(dx)$, $\Gamma(dy)$, and $\Gamma(d)$ without distinguishing among different execution paths, which are irrelevant in this case.

Generating Bound Evidence. Although we have shown how to generate constraints for Γ according to data dependence relations on \mathcal{G} , the presence of under-constrained variables in the analysis can still lead to imprecision, as illustrated in § 1. The key challenge is to infer additional information for these under-constrained variables when solely relying on data dependence relations gives imprecise results.

As motivated in § 1, a bound checking statement $\text{if } (v_1 \text{ cmp } v_2)$ reveals programmer’s beliefs about the values of the checked variables: a comparison between v_1 and v_2 suggests that the programmer may believe that $v_1 = v_2$, which we call a *bound evidence*. Bound evidence can aid in providing additional information for the under-constrained variables, e.g., it is used to infer $0 \in \Gamma(dz)$ as discussed in § 3. Therefore, we propose to generate additional constraints on

Γ for bound evidence. Specifically, Rule *bound-evi* of Figure 5 will unify the guarded symbolic value sets for the compared variables v_1 and v_2 to $\Gamma(v_1) \dot{\cup} \Gamma(v_2)$. These additional evidence constraints, together with the data dependence constraints generated before will be resolved later to propagate evidence according to the data dependence relations.

Example 4.3. In our motivating example of Figure 1, the two bound checking statements at Line 15 and Line 21 will cause the analysis to generate bound evidence. Applying Rule *bound-evi*, we obtain the additional constraints:

$$\Gamma(dx), \Gamma(dy) \leftarrow \Gamma(dx) \dot{\cup} \Gamma(dy) \quad (6)$$

$$\Gamma(dx), \Gamma(d) \leftarrow \Gamma(dx) \dot{\cup} \Gamma(d) \quad (7)$$

$$\Gamma(dy), \Gamma(d) \leftarrow \Gamma(dy) \dot{\cup} \Gamma(d) \quad (8)$$

We know from the above constraints that the symbolic values for dx , dy , and d are unified. As we will see later, this information refines the constraints generated in Example 4.2 and helps us to infer evidence for dz to be zero.

On-demand Constraints Generation. In Algorithm 1, the procedure *genConstraints* generates a system of constraints $\text{Cons}(\Gamma)$ on-demand from \mathcal{G} and a given node n . Specifically, we perform a backward depth-first traversal on \mathcal{G} starting from n (Line 9). This effectively computes a slice of \mathcal{G} affecting the value of n . Following the reverse node order of this DFS traversal, we collect constraints using rules in Figure 5 (Lines 10-12) and finally return these constraints at Line 13. The order ensures that rules for v_1 come before v_2 if v_1 is used to define v_2 ’s value, thus speeding up the convergence of the solving process introduced later, similar to data flow analysis [2].

4.2.2 Constraint Resolution. Given a set of constraints $\text{Cons}(\Gamma)$, the procedure *solve* in Algorithm 1 computes its solution, which is a mapping Γ from nodes to guarded symbolic value sets. At a high level, Γ starts with an initial state mapping any node to the empty set (Line 15) and is iteratively updated according to the constraints $\text{Cons}(\Gamma)$. Notice that each constraint in $\text{Cons}(\Gamma)$ is built from rules in Figure 5 and, thus, has the syntactic form $\Gamma(o) \leftarrow S$. For updating $\Gamma(o)$, we calculate the result of the set operation encoded by S , denoted by $\text{eval}(S, \Gamma)$ at Line 18. The calculation result is combined with the old value $\Gamma(o)$ using the $\dot{\cup}$ operator (c.f. Definition 4.3) to produce the updated value set S' for o (Lines 18-19), since \leftarrow represents a weak update. The algorithm finishes when no constraint causes an update to Γ anymore (Lines 16-21).

Example 4.4. For the constraints (4)-(5) shown in Example 4.2, Lines 18-19 of the *solve* procedure compute the guarded symbolic value set for dz as follows:

$$\Gamma(dz) = \{(\hat{dz}, \text{true})\} \dot{\cup}$$

$$\{(2, \text{true})\} \tilde{\wedge} \{(\hat{dx}, \text{true})\} \tilde{\wedge} \{(\hat{dy}, \text{true})\} \tilde{\wedge} \{(\hat{d}, \text{true})\}$$

$$= \{(2 \times \hat{dx} - (\hat{dy} + \hat{d}), \text{true}), (\hat{dz}, \text{true})\}$$

Constraints (6)-(8) from Example 4.3 further refines the result to:

$$\Gamma(dz) = \{(2 \times \hat{a} - (\hat{b} + \hat{c}), \text{true}), (\hat{dz}, \text{true}) \mid a, b, c \in \{dx, dy, d\}\}$$

Since $(0, \text{true}) \in \Gamma(dz)$, we have found evidence for the variable dz to trigger divide-by-zero.

Algorithm 2: Divide-by-zero bug detection by WIT.

Input: A program P .
Input: A divisor variable v and the location l for the division instruction.
Output: Whether divide-by-zero may happen at l

```

1 def decideBug( $P, v, l$ ):
2    $S \leftarrow \text{symbolicAnalysis}(P, v)$ 
3    $C \leftarrow$  path condition for reaching  $l$ 
4   foreach  $(va, cond) \in S$  do
5      $pc \leftarrow C \wedge cond$ 
6     if  $SAT(pc \wedge va = 0)$  then
7       if  $tainted(va)$  then
8         return true
9       else if  $UNSAT(pc \wedge va \neq 0)$  then
10        return true
11  end
12  return false

```

The above example shows the advantage of generating and resolving constraints from \mathcal{G} according to the data dependence relations. Our constraints are compact, i.e., resolving the constraints does not require reasoning about irrelevant control flow paths. In contrast, an exhaustive approach such as symbolic execution needs to enumerate all eight paths in the function `move` of Figure 1, just to figure out the guarded symbolic value set for dz :

Path	dz
$c_1 = dx \geq dy \wedge dx \geq d \wedge dy \geq d$	$2 \times dx - (dy + d)$
...	...
$c_8 = dx < dy \wedge dx < d \wedge dy < d$	$2 \times dx - (dy + d)$

The constraints built by exhaustive path enumeration would have been overly redundant:

$$\Gamma(dz) \leftrightarrow \text{trans}([(\Gamma(2) \widetilde{\times} \Gamma(dx) \widetilde{=} (\Gamma(dy) \widetilde{+} \Gamma(d)), c_1)])$$

...

$$\Gamma(dz) \leftrightarrow \text{trans}([(\Gamma(2) \widetilde{\times} \Gamma(dx) \widetilde{=} (\Gamma(dy) \widetilde{+} \Gamma(d)), c_8)])$$

Since $c_1 \vee \dots \vee c_8 \equiv \text{true}$, the above constraints are equivalent to our generated constraint (4) in Example 4.2.

One noteworthy point about the *solve* procedure is that it can produce a value set S' with a large size. To keep the analysis tractable, we join the additional values for variable v to its sound abstraction \hat{v} when the size of $\Gamma(v)$ has reached a predefined threshold. In this work, we choose the threshold to be 20 by experience.

4.3 Divide-by-Zero Bug Detection

Our system, WIT, for precise divide-by-zero detection is shown in Algorithm 2. Specifically, given a program P , a divisor variable v , and the location l of the division instruction, Algorithm 2 returns whether a divide-by-zero bug may happen at l . We first call *symbolicAnalysis* of Algorithm 1 using P and v as arguments to obtain the guarded symbolic value set S for v (Line 2). Divide-by-zero detection is achieved by enumerating each guarded value $(va, cond)$ in S to check for the bug condition.

To report a divide-by-zero bug, we adhere to the criteria $C1^* - C2^*$ mentioned in § 1. First, according to $C1^*$, the conjunction of the path

$$\text{taint-evi} \frac{v \in V_t, v \xrightarrow{cond} v', SAT(cond)}{tainted(v') = \text{true}}$$

Figure 6: Criteria for propagating the tainted flag.

condition and error condition should be satisfiable. In Algorithm 2, the path condition is denoted by pc , consisting of the condition C for reaching the division instruction and the guarding condition $cond$ for the checked value va (Line 5). The error condition is simply $va = 0$. Therefore, Line 6 checks the satisfiability of $pc \wedge va = 0$ for determining $C1^*$. Second, according to $C2^*$, the concerned symbolic value va should have affirmative evidence to be zero:

- **Source evidence:** If va comes from tainted input, we can report a divide-by-zero bug (Lines 7-8). For taint detection, we can apply any existing taint analysis to compute tainted nodes on \mathcal{G} . Specifically, we use Rule *taint-evi* in Figure 6: a node is tainted if it is data dependent on the tainted input and the condition of the value flow is satisfiable.
- **Bound evidence:** Otherwise, we report a divide-by-zero bug if va must be zero under the condition pc (Lines 9-10). This can happen for two reasons: either the symbolic value va is literally constant zero, or it is forced to be zero by the condition pc (e.g., when $pc = -1 < va < 1$). Recall that bound evidence is encoded as additional constraints in $Cons(\Gamma)$ during the analysis and can implicitly enforce va to be zero, it is thus more likely to satisfy the must query.

For the program in Figure 1, Algorithm 1 deduces that $(0, \text{true}) \in \Gamma(dz)$ as illustrated by Example 4.4. Thus, Lines 9-10 of Algorithm 2 report a divide-by-zero bug for the variable dz .

5 EVALUATION

We implement WIT based on LLVM [20]. Similar to previous works on static bug detection [3, 29, 30, 32], we unroll each loop once in control flow graphs and call graphs. Our experimental evaluation is designed to answer the following research questions:

- **RQ3:** The effectiveness of WIT. Particularly,
 - **RQ3.1:** How does our intuition of evidence-based reasoning affect the analysis precision?
 - **RQ3.2:** Compared to the conventional evidence agnostic method, how many bugs are missed by WIT? What are the reasons for missing them?
 - **RQ3.3:** Can WIT detect real-world divide-by-zero bugs?
- **RQ4:** How does WIT perform compared with existing divide-by-zero detectors?

5.1 Experimental Setup

Subjects. We have selected 12 open source C/C++ projects to perform the evaluation, shown in Table 2. Our selection criteria are as follows:

- **Popularity:** The selected projects are popular (e.g., have at least 10K stars on GitHub) and actively maintained.
- **Generality:** The projects cover different sizes (ranging from tens of thousands to tens of millions of lines of code) and functionalities (including operating system, image processing, database system, network library, etc.)

Table 2: Selected projects for evaluation.

Project	Loc	#Div/KLoC
masscan	34k	5.4
goaccess	53k	1.1
libuv	59k	0.8
redis	131k	5.0
git	226k	4.5
vim	354k	1.8
ImageMagick	382k	6.6
openssl	465k	4.1
systemd	600k	5.0
php	1,012k	1.3
gdb	1,932k	1.6
Linux kernel	15,164k	2.1

Table 3: Divide-by-zero detection on real-world projects. WIT^- represents a variant of WIT unaware of evidence.

Project	# of reports		FP rate		Analysis time	
	WIT	WIT^-	WIT	WIT^-	WIT	WIT^-
masscan	3	10	30%	80%	4m24s	6m
goaccess	2	15	0	80%	1m38s	1m33s
libuv	1	3	0	67%	1m32s	1m29s
redis	1	20	0	95%	23m12s	23m9s
git	10	29	40%	79%	38m3s	34m34s
vim	4	32	25%	88%	109m55s	109m12s
ImageMagick	5	47	20%	89%	196m39s	189m10s
openssl	2	18	0	89%	31m55s	31m48s
systemd	5	20	60%	90%	202m8s	183m21s
php	4	17	50%	88%	104m38s	94m21s
gdb	5	213	20%	96%	248m28s	299m1s
Linux kernel	53	2839	19%	NA	425m35s	452m12s

- Intensive use of division instruction: The projects perform division operation intensively, e.g., they contain 3.3 division instructions every 1k lines of code on average.

Environment. The experiments were performed on a computer with two 20 core processors Intel(R) Xeon(R) CPU E5-2698 v4@2.20GHz and 256GB physical memory running Ubuntu-16.04.

Open Data. The results of our empirical study and confirmed bugs detected by WIT are available at the link: <https://github.com/yiyuaner/ICSE-2022-Wit-data>.

5.2 Effectiveness of Divide-by-Zero Detection

To study the effectiveness of the evidence-based symbolic analysis algorithm, we compare WIT with its variant WIT^- . WIT^- is the evidence agnostic path-sensitive method. In WIT^- , we remove Rule *bound-evi* from Figure 5 and report an error whenever the condition is satisfiable, i.e., whenever the check at Line 6 of Algorithm 2 passes. All other aspects of WIT^- are the same as WIT .

Table 3 shows the experiment result. We compare the total number of bug reports, the false positive rate, and the analysis time for WIT and WIT^- . For each variant on a specific project, we examine its list of output reports and stop if the number of false positives has exceeded 200 (when this happens, we use “NA” to denote its false positive rate). WIT effectively reports 95 divide-by-zero bugs

Table 4: Distribution of true positives reported by WIT (the column “Total”) into Class Src and Class Bd. The column “Missed” shows the number of true positives reported by WIT^- but missed by WIT .

Project	Total	Src	Bd	Missed
masscan	2	2	0	0
goaccess	2	2	0	1
libuv	1	0	1	0
redis	1	1	0	0
git	6	4	2	0
vim	3	1	2	1
ImageMagick	4	3	2	1
openssl	2	1	1	0
systemd	2	1	1	0
php	2	0	2	0
gdb	4	3	1	4
Linux kernel	43	36	10	NA

Table 5: Divide-by-zero bugs confirmed by developers.

git	Linux	gdb	Image Magick	goaccess	libuv	openssl	vim	systemd
1	4	1	2	2	1	1	1	1

```

1 int blk_mq_map_queues(blk_mq_queue_map *qmap) {
2     unsigned *map = qmap->mq_map;
3     unsigned nr_queues = qmap->nr_queues;
4     unsigned cpu, first_sibling;
5     unsigned q = 0;
6     for_each_possible_cpu(cpu) {
7         if (q < nr_queues) {
8             map[cpu] = queue_index(qmap, nr_queues, q++);
9         } else {
10            if (first_sibling == cpu)
11                map[cpu] = queue_index(qmap, nr_queues, q++);
12            ...
13        }
14    }
15 }
16 int queue_index(blk_mq_queue_map *qmap,
17                unsigned nr_queues, const int q)
18 {
19     return qmap->queue_offset + (q % nr_queues);
20 }

```

Figure 7: A confirmed divide-by-zero bug in Linux kernel.

(72 of them are true positives) among 12 real-world applications, with a low average false positive rate of 22%. In contrast, WIT^- has an average false positive rate of 86%.

Answer to RQ3.1: The precision of WIT greatly outperforms its no-evidence counterpart, proving the significance of affirmative evidence for precise divide-by-zero detection.

Table 4 shows the distribution of true positives reported by WIT and the number of true positives detected by WIT^- but missed by WIT . From the table, we conclude that both source and bound evidence are useful in precisely detecting divide-by-zero bugs. Besides, WIT misses 12% of the true positives detected by WIT^- on average,

which is relatively low considering its drastic precision improvement. Most of the missed bugs are due to the incomplete modeling of taint sources, i.e., the analysis fails to identify certain APIs as taint sources, thus missing source evidence and fails to report certain bugs. WIT currently models some common taint sources from the standard C library but each target program may contain its specific taint sources that we do not know beforehand.

Answer to RQ3.2: WIT only misses a small proportion of bugs compared with its no-evidence counterpart WIT⁻, mainly due to its incomplete modeling of taint sources.

We report the detected true positives to developers and 14 bugs have been confirmed by the developers, as shown in Table 5. Figure 7 shows a real divide-by-zero bug WIT discovered from the Linux kernel. The function `queue_index` uses `nr_queues` as a divisor, which comes from the unknown input argument `qmap->nr_queues`. Since the analysis generates a bound evidence `q=nr_queues` from Line 7 of Figure 7 and `q` is assigned zero value in Line 5, it deduces that `nr_queues` may be zero and reports a divide-by-zero bug.

Answer to RQ3.3: WIT is capable of detecting real divide-by-zero bugs: it has found 72 divide-by-zero bugs, 14 of which have been confirmed by the developers.

5.3 Comparison with Existing Detectors

We compare WIT with three competing tools, namely Crab [18], Infer [5], and Clang Static Analyzer (CSA). Crab adopts numerical abstract interpretation, Infer takes the bi-abduction theorem proving technique, and CSA performs local symbolic execution.

Comparing with Crab. We instrument the program by asserting the divisor variable to be non-zero and use Crab to verify the inserted assertions with the interval domain [8]². When Crab fails to verify the safety of a division instruction, it emits a warning as a potential divide-by-zero bug. We cannot directly evaluate the precision of Crab because it currently does not provide an interface to map its verification result to source code location for reports examination. Instead, we record the number of warnings it generates and the ratio of its generated warnings over the total number of checked assertions. Intuitively, a tool reporting warnings too often (e.g., consider a tool that reports a bug for half of the division instructions in the program) maybe unusable in the industrial setting [4, 7]. However, it should be noted that this only loosely relates to the precision and our comparison has the caveat of possibly underestimating the effectiveness of Crab.

The analysis result for Crab is shown in Table 6. On average, Crab produces 212 bug reports per project, failing to verify 44% of the checked assertions on average. Such a large number of bug reports will impede the user from adopting the tool in divide-by-zero detection. In contrast, WIT reports 95 divide-by-zero bugs for the 12 projects in total, with an average false positive rate of 22%.

Comparing with Infer and CSA. Table 7 shows the results of Infer and CSA. From the data, we conclude that:

²We have tried more expressive domains. The zone domain [25] fails on 6 of the 12 projects due to time out or memory out in a 12 hours experiment, while the octagon domain [26] fails on 8 of the 12 projects.

Table 6: Analysis results for Crab.

Project	#Warnings	Warning ratio	Time
massscan	74	66%	1m43s
goaccess	9	53%	0m34s
libuv	6	27%	0m15s
redis	542	36%	200m30s
git	128	58%	175m49s
vim	174	42%	111m29s
ImageMagick	102	27%	389m49s
openssl	111	60%	5m26s
systemd	925	43%	68m41s
php	101	13%	141m30s
gdb	165	57%	20m10s
Linux kernel	OOM	OOM	OOM

Table 7: Divide-by-zero detection results for Infer and CSA. NA denotes the false positive rate when no bug is reported.

Project	# of reports		FP rate		Analysis time	
	Infer	CSA	Infer	CSA	Infer	CSA
massscan	0	0	NA	NA	1m4s	4m8s
goaccess	1	0	0	NA	3m27s	10m59s
libuv	0	0	NA	NA	3m13s	4m47s
redis	4	1	50%	0	8m47s	15m37s
git	7	0	100%	NA	8m55s	15min5s
vim	0	2	NA	50%	24m21s	17m47s
ImageMagick	7	2	71%	100%	16m7s	25m24s
openssl	0	1	NA	0	49m2s	9m43s
systemd	1	0	100%	NA	31m21s	11m45s
php	0	2	NA	100%	20m23s	57m28s
gdb	0	6	NA	83%	41m59s	81m17s
Linux kernel	Crash	63	Crash	51%	Crash	281m4s

- (1) Infer generates zero reports on six projects and fails on one project, while CSA generates zero reports on five projects. Thus, they both have missed many true positives detected by WIT. The relatively low recall is due to practical considerations for the tools, such as the limited capability for detecting cross-file bugs and inherent approximations in their algorithms.
- (2) Over the projects where the number of bug reports is non-zero, Infer has an average false positive rate of 64%, and CSA has an average false positive rate of 55%.
- (3) WIT runs slower than Infer and CSA in projects with large sizes. This is mainly because we utilize SMT solving to be fully path-sensitive. However, WIT finishes within 7.5 hours in all projects, which is an acceptable performance given its precision improvement.

Answer to RQ4: compared with existing divide-by-zero detectors, WIT is significantly more precise and sometimes even detects more divide-by-zero bugs.

6 RELATED WORK

Abstract Interpretation. In abstract interpretation [8], many abstract domains have been designed to verify numerical properties, such as the octagon [26] and polyhedra domain [9]. Recent works have tried to balance between precision and cost of abstract interpretation. Oh et al. [27] utilize an impact pre-analysis to apply context-sensitivity selectively, while LAIT [19] identifies and removes redundant constraints in numerical analysis to improve speed without hurting precision too much. Mansur et al. [24] propose to automatically tailor the configurations of abstract interpreters according to the code under analysis and resource constraints. Our work improves precision from a different angle by finding affirmative evidence for triggering the bug.

Symbolic Analysis for Numerical Bug Detection. To the best of our knowledge, we are the first to design a symbolic analysis algorithm targeting the problem of divide-by-zero detection. Previous works mainly focus on applying symbolic analysis in static buffer overflow detection and integer overflow detection. Li et al. [22] adopt a simple symbolic value representation and filter out irrelevant dependencies during the symbolic value computation. Marple [21] performs on-demand backward symbolic execution, starting from a buffer access statement and categorizing program paths to prioritize bug reports. SIFT [23] utilizes precondition inference in computing sound input filters to nullify integer overflow errors associated with critical program sites such as memory allocation or block copy sites. These works adopt path-sensitive symbolic domains similar to ours but do not discuss the imprecision problem brought by the under-constrained variables or how to address it, which is the key contribution of this work.

Reasoning About Programmer’s Beliefs. Engler et al. [16] first propose to infer from code about the programmers’ beliefs and detect potential bugs by spotting inconsistency of such beliefs. Dillig et al. [12] formalize the intuition of [16] in a framework for detecting semantic inconsistency, where they design a null pointer dereference checker. The belief propagation is carried out using type inference: e.g., a pointer inferred with a “possibly NULL” type should not be used in a context that requires a “not NULL” type. However, divide-by-zero detection requires an infinite analysis domain and existing approaches for type-state properties are not applicable.

Inspired by the success of utilizing beliefs in previous works on static analysis, we firstly introduce the use of beliefs to address the imprecision problem brought by under-constrained variables in divide-by-zero detection. We generalize the idea of “programmer’s belief” in analyzing numerical computation by generating bound evidence from bound checking statements in the program. Although we do not directly infer inconsistency, the evidence we generate serves as additional constraints, which greatly improves the overall precision of the analysis. The imprecision problem caused by under-constrained variables in static analysis is first observed in [13], where they tackle the problem with constraints rewriting. Since they do not consider numerical property in the constraints and assume a finite analysis domain, their approach is not applicable in our scenario.

7 CONCLUSION

We have proposed WIT, a static analysis method to find divide-by-zero bugs with affirmative evidence. The analysis looks for affirmative evidence, namely source evidence and bound evidence, to find divide-by-zero bugs with high confidence. It has been shown effective in detecting divide-by-zero bugs precisely in large-scale real-world software.

8 ACKNOWLEDGMENTS

We thank the anonymous reviewers for their insightful comments. The authors are supported by the RGC16206517, ITS/440/18FP and PRP/004/21FX grants from the Hong Kong Research Grant Council and the Innovation and Technology Commission, Ant Group, and the donations from Microsoft and Huawei. Peisen Yao is the corresponding author.

REFERENCES

- [1] [n.d.]. The Economics of Static Analysis Tool Usage. <https://blogs.grammotech.com/the-economics-static-analysis-tool-usage>. Online; accessed 12-August-2021.
- [2] Alfred V Aho, Monica S Lam, Ravi Sethi, and Jeffrey D Ullman. 2020. *Compilers: principles, techniques and tools*.
- [3] Domagoj Babic and Alan J. Hu. 2008. Calysto: Scalable and Precise Extended Static Checking. In *Proceedings of the 30th International Conference on Software Engineering (Leipzig, Germany) (ICSE '08)*. Association for Computing Machinery, New York, NY, USA, 211–220. <https://doi.org/10.1145/1368088.1368118>
- [4] Al Bessey, Ken Block, Ben Chelf, Andy Chou, Bryan Fulton, Seth Hallen, Charles Henri-Gros, Asya Kamsky, Scott McPeak, and Dawson Engler. 2010. A Few Billion Lines of Code Later: Using Static Analysis to Find Bugs in the Real World. *Commun. ACM* 53, 2 (Feb. 2010), 66–75. <https://doi.org/10.1145/1646353.1646374>
- [5] Cristiano Calcagno, Dino Distefano, Peter O’Hearn, and Hongseok Yang. 2009. Compositional Shape Analysis by Means of Bi-Abduction. In *Proceedings of the 36th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (Savannah, GA, USA) (POPL '09)*. Association for Computing Machinery, New York, NY, USA, 289–300. <https://doi.org/10.1145/1480881.1480917>
- [6] Sigmund Cherm, Lonnie Princehouse, and Radu Rugina. 2007. Practical Memory Leak Detection Using Guarded Value-Flow Analysis. In *Proceedings of the 28th ACM SIGPLAN Conference on Programming Language Design and Implementation (San Diego, California, USA) (PLDI '07)*. Association for Computing Machinery, New York, NY, USA, 480–491. <https://doi.org/10.1145/1250734.1250789>
- [7] Maria Christakis and Christian Bird. 2016. What Developers Want and Need from Program Analysis: An Empirical Study. In *Proceedings of the 31st IEEE/ACM International Conference on Automated Software Engineering (Singapore, Singapore) (ASE 2016)*. Association for Computing Machinery, New York, NY, USA, 332–343. <https://doi.org/10.1145/2970276.2970347>
- [8] Patrick Cousot and Radhia Cousot. 1977. Abstract Interpretation: A Unified Lattice Model for Static Analysis of Programs by Construction or Approximation of Fixpoints. In *Proceedings of the 4th ACM SIGACT-SIGPLAN Symposium on Principles of Programming Languages (Los Angeles, California) (POPL '77)*. Association for Computing Machinery, New York, NY, USA, 238–252. <https://doi.org/10.1145/512950.512973>
- [9] Patrick Cousot and Nicolas Halbwachs. 1978. Automatic Discovery of Linear Restraints among Variables of a Program. In *Proceedings of the 5th ACM SIGACT-SIGPLAN Symposium on Principles of Programming Languages (Tucson, Arizona) (POPL '78)*. Association for Computing Machinery, New York, NY, USA, 84–96. <https://doi.org/10.1145/512760.512770>
- [10] Ron Cytron, Jeanne Ferrante, Barry K Rosen, Mark N Wegman, and F Kenneth Zadeck. 1991. Efficiently computing static single assignment form and the control dependence graph. *ACM Transactions on Programming Languages and Systems (TOPLAS)* 13, 4 (1991), 451–490.
- [11] Barthélemy Dagenais and Laurie Hendren. 2008. Enabling Static Analysis for Partial Java Programs. In *Proceedings of the 23rd ACM SIGPLAN Conference on Object-Oriented Programming Systems Languages and Applications (Nashville, TN, USA) (OOPSLA '08)*. Association for Computing Machinery, New York, NY, USA, 313–328. <https://doi.org/10.1145/1449764.1449790>
- [12] Isil Dillig, Thomas Dillig, and Alex Aiken. 2007. Static Error Detection Using Semantic Inconsistency Inference. In *Proceedings of the 28th ACM SIGPLAN Conference on Programming Language Design and Implementation (San Diego, California, USA) (PLDI '07)*. Association for Computing Machinery, New York, NY, USA, 435–445. <https://doi.org/10.1145/1250734.1250784>

- [13] Isil Dillig, Thomas Dillig, and Alex Aiken. 2010. Reasoning about the Unknown in Static Analysis. *Commun. ACM* 53, 8 (Aug. 2010), 115–123. <https://doi.org/10.1145/1787234.1787259>
- [14] Isil Dillig, Thomas Dillig, and Alex Aiken. 2012. Automated Error Diagnosis Using Abductive Inference. *SIGPLAN Not.* 47, 6 (June 2012), 181–192. <https://doi.org/10.1145/2345156.2254087>
- [15] Dino Distefano, Manuel Fähndrich, Francesco Logozzo, and Peter W. O'Hearn. 2019. Scaling Static Analyses at Facebook. *Commun. ACM* 62, 8 (July 2019), 62–70. <https://doi.org/10.1145/3338112>
- [16] Dawson Engler, David Yu Chen, Seth Hallem, Andy Chou, and Benjamin Chelf. 2001. Bugs as Deviant Behavior: A General Approach to Inferring Errors in Systems Code. In *Proceedings of the Eighteenth ACM Symposium on Operating Systems Principles* (Banff, Alberta, Canada) (*SOSP '01*). Association for Computing Machinery, New York, NY, USA, 57–72. <https://doi.org/10.1145/502034.502041>
- [17] Jeanne Ferrante, Karl J. Ottenstein, and Joe D. Warren. 1987. The Program Dependence Graph and Its Use in Optimization. *ACM Trans. Program. Lang. Syst.* 9, 3 (July 1987), 319–349. <https://doi.org/10.1145/24039.24041>
- [18] Arie Gurfinkel, Temesghen Kahsai, Anvesh Komuravelli, and Jorge A. Navas. 2015. The SeaHorn Verification Framework. In *Computer Aided Verification*, Daniel Kroening and Corina S. Păsăreanu (Eds.). Springer International Publishing, Cham, 343–361.
- [19] Jingxuan He, Gagandeep Singh, Markus Püschel, and Martin Vechev. 2020. Learning Fast and Precise Numerical Analysis. In *Proceedings of the 41st ACM SIGPLAN Conference on Programming Language Design and Implementation* (London, UK) (*PLDI 2020*). Association for Computing Machinery, New York, NY, USA, 1112–1127. <https://doi.org/10.1145/3385412.3386016>
- [20] C. Lattner and V. Adve. 2004. LLVM: a compilation framework for lifelong program analysis and transformation. In *International Symposium on Code Generation and Optimization, 2004. CGO 2004*. 75–86. <https://doi.org/10.1109/CGO.2004.1281665>
- [21] Wei Le and Mary Lou Soffa. 2008. Marple: A Demand-Driven Path-Sensitive Buffer Overflow Detector. In *Proceedings of the 16th ACM SIGSOFT International Symposium on Foundations of Software Engineering* (Atlanta, Georgia) (*SIGSOFT '08/FSE-16*). Association for Computing Machinery, New York, NY, USA, 272–282. <https://doi.org/10.1145/1453101.1453137>
- [22] Lian Li, Cristina Cifuentes, and Nathan Keynes. 2010. Practical and Effective Symbolic Analysis for Buffer Overflow Detection. In *Proceedings of the Eighteenth ACM SIGSOFT International Symposium on Foundations of Software Engineering* (Santa Fe, New Mexico, USA) (*FSE '10*). Association for Computing Machinery, New York, NY, USA, 317–326. <https://doi.org/10.1145/1882291.1882338>
- [23] Fan Long, Stelios Sidiropoulos, Deokhwan Kim, and Martin Rinard. 2014. Sound Input Filter Generation for Integer Overflow Errors. In *Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages* (San Diego, California, USA) (*POPL '14*). Association for Computing Machinery, New York, NY, USA, 439–452. <https://doi.org/10.1145/2535838.2535888>
- [24] Muhammad Numair Mansur, Benjamin Mariano, Maria Christakis, Jorge A Navas, and Valentin Wüstholtz. 2021. Automatically Tailoring Abstract Interpretation to Custom Usage Scenarios. In *International Conference on Computer Aided Verification*. Springer, 777–800.
- [25] Antoine Miné. 2001. A new numerical abstract domain based on difference-bound matrices. In *Symposium on Program as Data Objects*. Springer, 155–172.
- [26] Antoine Miné. 2006. The Octagon Abstract Domain. *Higher Order Symbol. Comput.* 19, 1 (March 2006), 31–100. <https://doi.org/10.1007/s10990-006-8609-1>
- [27] Hakjoo Oh, Wonchan Lee, Kihong Heo, Hongseok Yang, and Kwangkeun Yi. 2014. Selective Context-Sensitivity Guided by Impact Pre-Analysis. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation* (Edinburgh, United Kingdom) (*PLDI '14*). Association for Computing Machinery, New York, NY, USA, 475–484. <https://doi.org/10.1145/2594291.2594318>
- [28] David A. Ramos and Dawson Engler. 2015. Under-Constrained Symbolic Execution: Correctness Checking for Real Code. In *24th USENIX Security Symposium (USENIX Security 15)*. USENIX Association, Washington, D.C., 49–64. <https://www.usenix.org/conference/usenixsecurity15/technical-sessions/presentation/ramos>
- [29] Qingkai Shi, Xiao Xiao, Rongxin Wu, Jinguo Zhou, Gang Fan, and Charles Zhang. 2018. Pinpoint: Fast and Precise Sparse Value Flow Analysis for Million Lines of Code. In *Proceedings of the 39th ACM SIGPLAN Conference on Programming Language Design and Implementation* (Philadelphia, PA, USA) (*PLDI 2018*). Association for Computing Machinery, New York, NY, USA, 693–706. <https://doi.org/10.1145/3192366.3192418>
- [30] Yulei Sui, Ding Ye, and Jingling Xue. 2012. Static Memory Leak Detection Using Full-Sparse Value-Flow Analysis. In *Proceedings of the 2012 International Symposium on Software Testing and Analysis* (Minneapolis, MN, USA) (*ISSTA 2012*). Association for Computing Machinery, New York, NY, USA, 254–264. <https://doi.org/10.1145/2338965.2336784>
- [31] Omer Tripp, Marco Pistoi, Stephen J. Fink, Manu Sridharan, and Omri Weisman. 2009. TAJ: Effective Taint Analysis of Web Applications. In *Proceedings of the 30th ACM SIGPLAN Conference on Programming Language Design and Implementation* (Dublin, Ireland) (*PLDI '09*). Association for Computing Machinery, New York, NY, USA, 87–97. <https://doi.org/10.1145/1542476.1542486>
- [32] Yichen Xie and Alex Aiken. 2005. Scalable Error Detection Using Boolean Satisfiability. In *Proceedings of the 32nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages* (Long Beach, California, USA) (*POPL '05*). Association for Computing Machinery, New York, NY, USA, 351–363. <https://doi.org/10.1145/1040305.1040334>